

Vector Database Indexing Benchmark

Weaviate HNSW vs Exact Baseline

This report presents a comparative analysis of approximate and exact vector indexing strategies under realistic dataset scale.

- Dataset: Mus�umPillerum / General-Knowledge

Dataset Description

- Query count: ~1,000 questions
- Domain: Open-domain general knowledge


Empty and invalid entries were filtered during preprocessing to ensure robust embedding generation.

1. Exact Baseline Index
 - Brute-force cosine similarity
 - Provides upper bound on recall
 - Computationally expensive


Indexing Strategies Evaluated

2. HNSW (Weaviate)
 - Approximate nearest neighbour graph
 - Tunable parameters: ef, efConstruction, maxConnections
 - Optimised for low-latency retrieval

Query Latency Across Indexing Strategies



Recall Comparison Across Indexes



- Exact search achieves maximum recall at significantly higher latency.
 - HNSW configurations offer substantial latency reductions with
- Observations and Trade-offs**
- Parameter tuning enables balanced trade-offs depending on workload.

These results demonstrate the necessity of approximate indexing strategies for scalable LLM-backed retrieval systems.